

# CIS 345/545    Project 3    Spring 2020

(Due April 27)

In this project, you will implement a tiny file system on top of a virtual disk. This file system is called TinyFS which is similar to the Unix file system discussed in the class notes. This virtual disk is actually a single 4MB file that is stored on the real file system in our lab.

To uncompress and extract the files to your working directory, login a Linux workstation in FH133E and then type


```
tar xvfz ~cis345s/pub/tinyfs.tar.gz
```

## Overview

The TinyFS has three major components: the shell, the file system itself, and the emulated disk, as shown to the right. Your task is to implement the middle component: the file system.

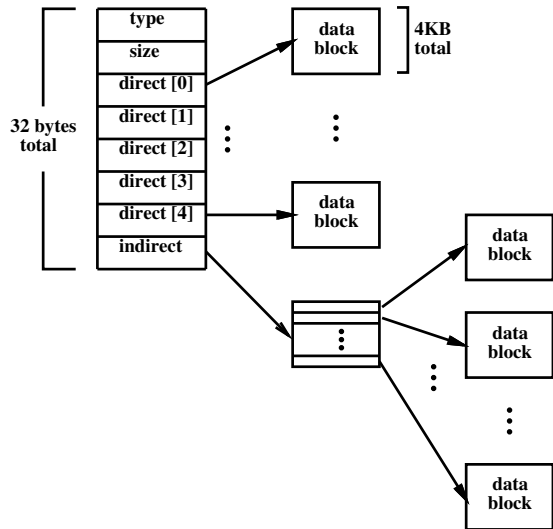
At the top level a user gives typed commands to a shell, instructing it to copy data out of the file system, delete a file, etc. The shell converts these typed commands into high-level operations on the file system, such as `tfs_read`, `tfs_getsize`, and `tfs_delete`. The file system is responsible for accepting these operations on files and converting them into simple block read and write operations on the emulated disk, called `disk_read`, and `disk_write`. The emulated disk, in turn, stores all of its data in an image file in the file system.

TinyFS has the following layout on disk. It assumes that the disk blocks are the common size of 4KB. The first block of the disk (i.e. disk block zero) is a "superblock" that describes the layout of the rest of the file system. The block following the superblock contains inode data structures. The remaining blocks in the file system are used as data blocks which contain either file data, directory entries, or occasionally indirect pointers. Below is a picture of a TinyFS image:



The superblock describes the layout of the rest of the file system. Each of the first four fields of the superblock is a 4-byte (32-bit) integer. The first field `signature` is always the magic number `TFS_MAGIC` (`0xc3450545`). When the file system is being used, the OS looks for this signature value. If it is correct, then the disk is assumed to contain a valid file system.

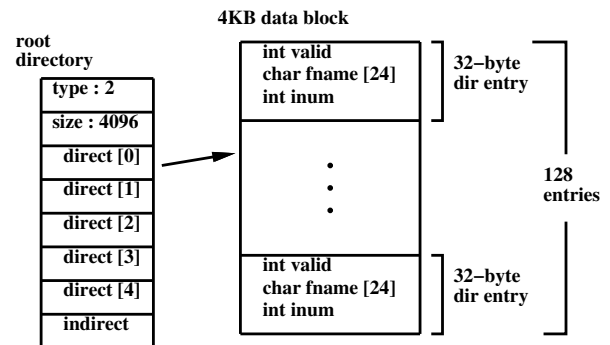
The remaining fields in the superblock describe the layout of the file system. `int num_blocks` is the total number of blocks on the disk. To make things easier, we use a fixed disk size 4MB in this project. So the total number of blocks is always 1024. We also assume the total number of inodes `num_inodes` is 128. Furthermore, `root_inode` contains the inode number (i.e. 1 in this project) of the root directory. Note that the inode number 0 is reserved as `NULL`. The arrays `block_in_use` and `inode_in_use` are the bitmaps of the blocks and inodes being used, respectively. The remainder of the superblock is left unused.



The left picture shows an inode entry in TinyFS. Each field of the inode is a 4-byte (32-bit) integer. The `type` field is 1 (i.e. the pre-defined constant `REGULAR`) if the inode is used for a regular file and is 2 (i.e. the pre-defined constant `DIR`) if it is used for a directory. Note that all files in TinyFS are stored in a single root directory. Hence, only the root inode has the `type` set to be 2. The `size` field contains the logical size of the data in bytes. For the root inode, its `size` is set to be one data block size 4096 (in this project). There are 5 `direct` pointers to data blocks, and one pointer to an indirect data block. In this context, "pointer" simply means the number of a block where data may be found. A value of zero may be used to indicate a `NULL` block pointer. Each inode occupies 32 bytes, so there are 128 inodes in each 4KB inode block. Note that an indirect data block is just a big

array of pointers to further data blocks. Each pointer is a 4-byte integer, and each block is 4KB, so there are 1024 pointers per block.

The right figure shows the structure of the root inode and its directory entries. Because the file system is very small, the root directory only needs one data block to store directory entries. The size of each directory entry is 32-byte and there are 128 entries in a block. There are three fields in a directory entry. The `valid` is 1 if the entry is valid (i.e. being used) and is 0 otherwise. For a valid entry, the `fname` field contains the file name and the `inum` contains the file's corresponding inode number.



## Disk Emulator

We provide you with a disk emulator on which to store your file system. As mentioned before, this virtual disk is actually stored as one big file in the file system, so that you can retrieve data in a "disk image". In addition, we will provide you with some sample disk images that you can experiment with to test your file system. Just like a real disk, the emulator only allows operations on entire disk blocks of 4KB (`DISK_BLOCK_SIZE`). You cannot read or write any smaller unit than that. The primary challenge of building a file system is converting the user's requested operations on arbitrary amounts of data into operations on fixed block sizes.

The interface to the simulated disk is given in `disk.h`:

```
#define DISK_BLOCK_SIZE 4096
int disk_init( const char *filename, int nblocks );
void disk_read( int blocknum, char *data );
void disk_write( int blocknum, const char *data );
void disk_close();
```

Before performing any sort of operation on the disk, you must call `disk_init` and specify a (real) disk image for storing the disk data, and the number of blocks (i.e. 1024 in this project) in the simulated disk. If this function is called on a disk image that already exists, the contained data will not be changed. When you are done using the disk, call `disk_close` to release the file. These two calls are already made for you in the

shell, so you should not have to change them. Furthermore, as the names suggest, the functions `disk_read` and `disk_write` read and write one block of data on the disk. Notice that the first argument is a block number, so a call to `disk_read(0,data)` reads the first 4KB of data on the disk, and `disk_read(1,data)` reads the next 4KB block of data on the disk. Every time that you invoke a read or a write, you must ensure that data points to a full 4KB of memory.

## File System

Using the existing simulated disk, you will build a working file system. Take note that we have already constructed the interface to the file system and provided some skeleton code. The interface is given in `fs.h`:

```
void tfs_debug();
int  tfs_delete(const char * );
int  tfs_getsize(const char *);
int  tfs_get_inumber(const char *);
int  tfs_read( int inumber, char *data, int length, int offset );
```

The various functions must work as follows:

- `tfs_debug` - This function scans a file system and reports on how the inodes and blocks are organized. Once you are able to scan and report upon the file system structures, the rest is easy. Your output from `tfs_debug` should be similar to the following:

```
superblock:
    signature is valid
    26 blocks in use
    4 inodes in use
root inode 1:
    size: 4096 bytes
    direct blocks: 2
foo inode 2:
    size: 13 bytes
    direct blocks: 3
bar inode 6:
    size 81929 bytes
    direct blocks: 10 11 15 17 18
    indirect block: 23
    indirect data blocks: 30 39 42 ...
```

- `tfs_delete` - It deletes the file specified by the parameter. Release all data and indirect blocks assigned to this inode and return them to the free block map. On success, return (positive) inode number. On failure, return 0.
- `tfs_getsize` - It returns the logical size of the given file, in bytes. Note that zero is a valid logical size for a file. On failure, return -1.
- `tfs_get_inumber` - This function searches the root directory entries and returns the inode number of the given file. On failure, return 0.
- `tfs_read` - This function reads data from the specified inode number `inumber`. That is, the user who calls this function needs to get the inode number via `tfs_get_inumber` first. This function copies "length" bytes from the inode into the "data" pointer, starting at "offset" in the inode. The total

number of bytes read will be returned. The number of bytes actually read could be smaller than the number of bytes requested, perhaps if the end of the file is reached. If the given `inumber` is invalid, or any other error is encountered, return 0.

It's quite likely that the file system module will need a number of global variables in order to keep track of the currently file system. For example, you will certainly need a global variable to keep track of the current free block bitmap, and perhaps other items as well. In the context of operating systems, they are common and quite normal. Furthermore, the file system has to maintain state persistently. You need to write all modifications, such as the inode block, superblock, etc., back to disk.

## Shell Interface

After you extract the source code, build it with `make`. We have provided for you a simple shell that will be used to exercise your filesystem and the simulated disk. When grading your work, we will use the shell to test your code, so be sure to test extensively. To use the shell, simply run `tinyfs` with the name of a provided disk image (i.e. `image0`, `image1`, or `image2`. For example, to use the `image0` example given below, run:

```
./tinyfs image0
```

Once the shell starts, you can use the `help` command to list the available commands:

```
TinyFS> help
Commands are:
  debug
  delete filename
  cat filename
  copyout <src_tfsfile> <dstfile>
  help
  quit
  exit
TinyFS>
```

Most of the commands correspond closely to the filesystem interface. For example, `debug`, and `delete` call the corresponding functions in the filesystem. The command `cat` reads an entire file out of the filesystem and displays it on the console, just like the Unix command of the same name. The command `copyout` copies a file (i.e. `<src_tfsfile>`) from the your emulated filesystem into the local Unix filesystem (i.e. `<dstfile>`).

## Implementation Notes

Your job is to implement TinyFS as described above by filling in the implementation of the file `fs.c`. You are not allowed to change any other code modules. We have already created some data structures to get you started. These can be found in `fs.c`.

Note that a raw 4 KB disk block can be used to represent five different kinds of data: a superblock, a block of 128 inodes, a directory of 128 entries, an indirect pointer block, or a plain data block. This presents a bit of a software engineering problem: how do we transform the raw data returned by `disk_read` into each of the five kinds of data blocks?

C provides a nice bit of syntax for exactly this problem. We can declare a union of each of our five different data types. A union looks like a struct, but forces all of its elements to share the same memory space. You can think of a union as several different types, all overlaid on top of each other.

```

union tfs_block {
    struct tfs_superblock super;
    struct tfs_inode inode[INODES_PER_BLOCK];
    struct tfs_dir_entry dentry[NUM_DENTRIES_PER_BLOCK];
    int pointers[POINTERS_PER_BLOCK];
    char data[DISK_BLOCK_SIZE];
};

```

Note that the size of a `tfs_block` union will be exactly 4KB: the size of the largest members of the union. To declare a variable of type:

```
union tfs_block block;
```

Now, we may use `disk_read` to load in the raw data from block zero. We give `disk_read` the variable `block.data`, which looks like an array of characters:

```
disk_read(0, block.data);
```

But, we may interpret that data as if it were a struct superblock by accessing the `super` part of the union. For example, to extract the signature of the super block, we might do this:

```
x = block.super.signature;
```

To check the block `k` is being used or not, we can check its bitmap:

```
if(block.super.block_in_use[k/32] & (1 << (k%32)) )
```

To set the corresponding bitmap of block `k`, we can use

```
block.super.block_in_use[k/32] |= (1 << (k%32));
```

On the other hand, suppose that we wanted to load disk block 49, assume that it is an indirect block, and then examine the 3rd pointer. Again, we would use `disk_read` to load the raw data:

```
disk_read(49, block.data);
```

But then use the pointer part of the union like so:

```
x = block.pointers[3];
```

The union offers a convenient way of viewing the same data from multiple perspectives. When we load data from the disk, it is just a 4 KB raw chunk of data (`block.data`). But, once loaded, the file system layer knows that this data has some structure. The file system layer can view the same data from another perspective by choosing another field in the union.

## Turning it in

Each group (at most two students) has to submit your program and report electronically. Use the following command on `grail` (under the `TinyFS` directory):

```
make submit
```

The report file `report.pdf` which includes the description of your code, experiences in debugging/testing, etc. should be put under the `TinyFS` directory. The cover page should contain your picture(s) (taken in FH128 using iMac), name(s) and the login id you used to turnin the project. Start on time and good luck. If you have any questions, send e-mail to [sang@cis.csuohio.edu](mailto:sang@cis.csuohio.edu).